

Applied Machine Learning

Regression Model Evaluation on Sydney housing dataset

Yunjun Xia

Data Source: <https://www.kaggle.com/shree1992/housedata>
(<https://www.kaggle.com/shree1992/housedata>)

Part 1: Data Description

```
In [1]: import pandas as pd  
import numpy as np
```

```
In [2]: house = pd.read_csv('data.csv')  
house = house.drop(columns = 'date')
```

```
In [3]: house.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 4600 entries, 0 to 4599  
Data columns (total 17 columns):  
price                4600 non-null float64  
bedrooms             4600 non-null float64  
bathrooms            4600 non-null float64  
sqft_living          4600 non-null int64  
sqft_lot             4600 non-null int64  
floors               4600 non-null float64  
waterfront           4600 non-null int64  
view                 4600 non-null int64  
condition            4600 non-null int64  
sqft_above           4600 non-null int64  
sqft_basement        4600 non-null int64  
yr_built             4600 non-null int64  
yr_renovated         4600 non-null int64  
street               4600 non-null object  
city                 4600 non-null object  
statezip             4600 non-null object  
country              4600 non-null object  
dtypes: float64(4), int64(9), object(4)  
memory usage: 611.1+ KB
```

By looking at the information of the house sales dataset, we can find that features of type "object" are categorical features while other features with type "float" or "int" are continuous features. Categorical Features: street, city, statezip, country. Continuous Features: price, bedrooms, bathrooms, sqft_living, sqft_lot, floors, waterfront, view, condition, sqft_above, sqft_basement, yr_built, yr_renovated.

```
In [4]: index = []
        for i in range(len(list(house.price))):
            if (list(house.price)[i] == 0) or (list(house.price)[i] > 10000000):
                index.append(i)
        df_house = house.drop(index)
        df_house.shape
```

```
Out[4]: (4549, 17)
```

Note:

I drop the rows with sales price = 0 or sales price > 10^7 .

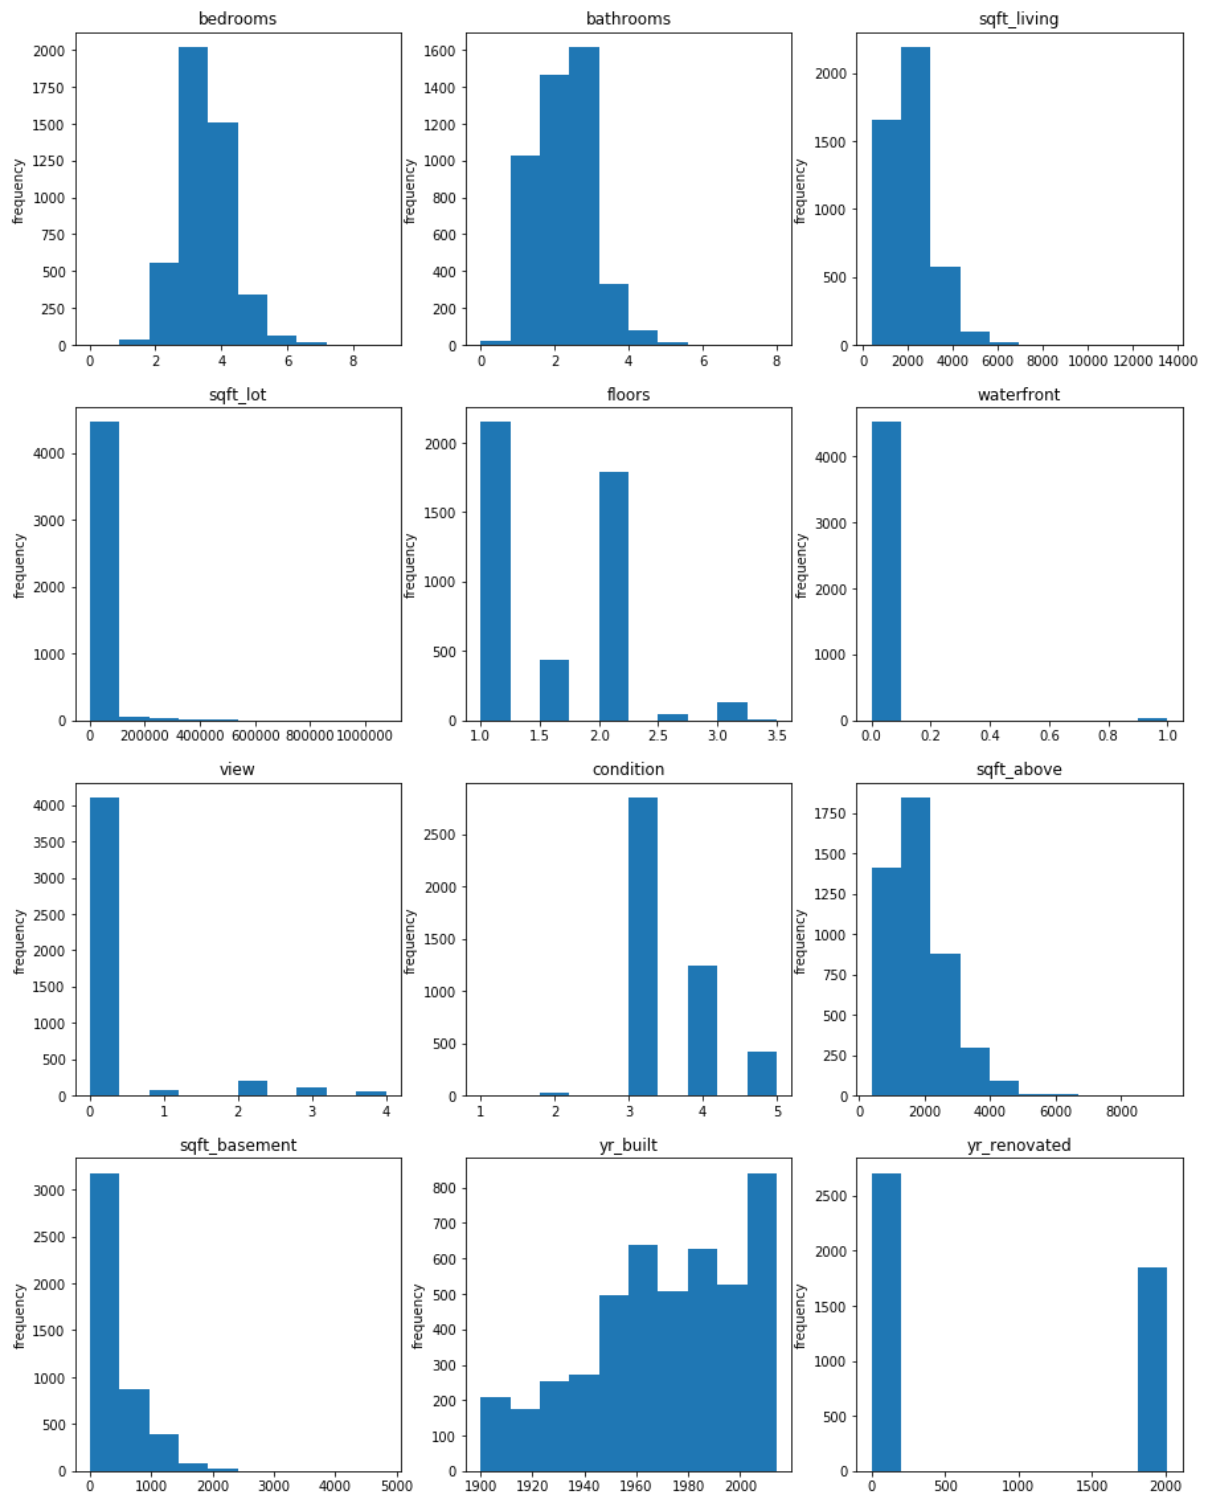
Part 2: Continuous Feature and Target Distribution

```
In [5]: import matplotlib.pyplot as plt
        %matplotlib inline
```

```

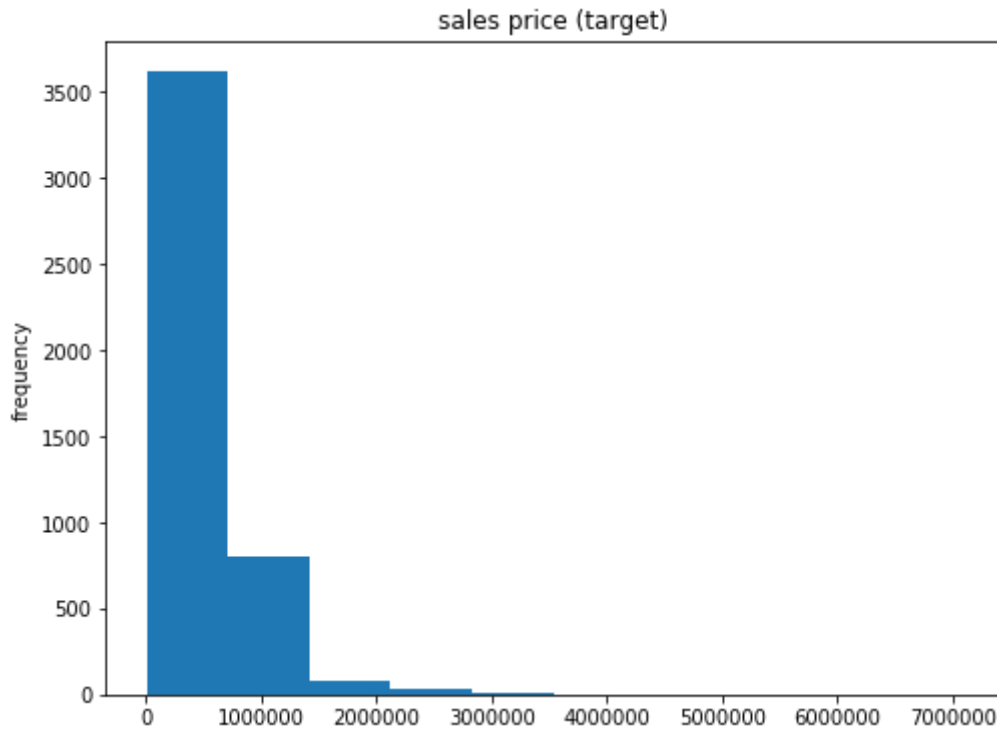
In [6]: fig, ax = plt.subplots(4, 3, figsize = (15,20))
count = 1
for i in range(4):
    for j in range(3):
        ax[i,j].hist(df_house.iloc[:,count])
        ax[i,j].set_title(str(df_house.columns[count]))
        ax[i,j].set_ylabel('frequency')
        count = count + 1

```



```
In [7]: plt.figure(figsize = (8, 6))  
plt.hist(df_house.price)  
plt.title('sales price (target)')  
plt.ylabel('frequency')
```

```
Out[7]: Text(0, 0.5, 'frequency')
```



One problem is that the dataset has outliers(values being too large) for many of the continuous features(like `sqft_lot`, `sqft_living`, `sqft_basement`) and the target(price). Thus, it's hard to visualize the distribution of price and continuous features without treatment to outliers. In this case, we can drop outliers to make a better visualization.

Another problem is that for the "yr_renovated" feature and the "sqft_basement" feature, there are so many data with value = 0. This can be viewed as missing values. It's also hard to visualize the distribution of this feature without treatment to the 0 values.

Moreover, since there are only two unique values for "waterfront", maybe we can treat this feature as a categorical feature.

Part 3: Dependency of target on each feature

```

In [8]: fig, ax = plt.subplots(4, 3, figsize = (15,20))
count = 1
for i in range(4):
    for j in range(3):
        ax[i,j].scatter(df_house.iloc[:,count], df_house.price)
        ax[i,j].set_title(str(df_house.columns[count]))
        ax[i,j].set_ylabel('price')
        ax[i,j].set_xlabel(str(df_house.columns[count]))
        count = count + 1

```



Part 4: Preprocessing and Model Fit: OLS, Ridge, Lasso and ElasticNet

```
In [9]: from sklearn.model_selection import train_test_split
        from sklearn.model_selection import cross_val_score
        from sklearn.preprocessing import StandardScaler, OneHotEncoder
        from sklearn.compose import make_column_transformer
        from sklearn.pipeline import make_pipeline
        from sklearn.linear_model import LinearRegression, Ridge, Lasso, Elastic
        Net
        from sklearn.impute import SimpleImputer
        import warnings
        warnings.filterwarnings('ignore')
```

```
In [10]: set(df_house.country)
```

```
Out[10]: {'USA'}
```

```
In [11]: len(set(df_house.street))
```

```
Out[11]: 4474
```

Note:

Since there is only one category "USA" for feature "country" and there are 4476 categories for feature "street", we drop these two columns for better fitting models. Also, let's declare the 0 values of "yr_renovated" and "sqft_basement" to be missing values.

```
In [12]: df_house_clean = df_house.drop(columns = 'country')
        df_house_clean = df_house_clean.drop(columns = 'street')
        df_house_clean.yr_renovated = df_house_clean.yr_renovated.replace(0, np.
        nan)
        df_house_clean.sqft_basement = df_house_clean.sqft_basement.replace(0, n
        p.nan)
```

```

In [13]: X_train, X_test, y_train, y_test = train_test_split(df_house_clean.iloc
[:,1:],
                                                    df_house_clean.iloc
[:,0],
                                                    random_state = 1)

categorical = df_house_clean.iloc[:,1:].dtypes == 'object'

conti_pre = make_pipeline(SimpleImputer(strategy = 'median'))
conti_pre_scaled = make_pipeline(SimpleImputer(strategy = 'median'), Sta
ndardScaler())

cat_pre = make_pipeline(SimpleImputer(strategy = 'constant', fill_value
= 'NA'),
                        OneHotEncoder(handle_unknown='ignore'))

preprocess = make_column_transformer((conti_pre, ~categorical),
                                     (cat_pre, categorical))

preprocess_scaled = make_column_transformer((conti_pre_scaled, ~categori
cal),
                                           (cat_pre, categorical))

```

```

In [14]: # Linear Regression
model_lr = make_pipeline(preprocess, LinearRegression())
model_lr_scaled = make_pipeline(preprocess_scaled, LinearRegression())

scores_lr = cross_val_score(model_lr, X_train, y_train)
scores_lr_scaled = cross_val_score(model_lr_scaled, X_train, y_train)

# Ridge
model_r = make_pipeline(preprocess, Ridge())
model_r_scaled = make_pipeline(preprocess_scaled, Ridge())

scores_r = cross_val_score(model_r, X_train, y_train)
scores_r_scaled = cross_val_score(model_r_scaled, X_train, y_train)

# Lasso
model_l = make_pipeline(preprocess, Lasso())
model_l_scaled = make_pipeline(preprocess_scaled, Lasso())

scores_l = cross_val_score(model_l, X_train, y_train)
scores_l_scaled = cross_val_score(model_l_scaled, X_train, y_train)

# ElasticNet
model_e = make_pipeline(preprocess, ElasticNet())
model_e_scaled = make_pipeline(preprocess_scaled, ElasticNet())

scores_e = cross_val_score(model_e, X_train, y_train)
scores_e_scaled = cross_val_score(model_e_scaled, X_train, y_train)

```

```
In [15]: print(f'The 5-fold cv mean score of linear regression is: {np.mean(scores_lr)}')
print(f'The 5-fold cv mean score of linear regression with scaling is: {np.mean(scores_lr_scaled)}\n')

print(f'The 5-fold cv mean score of ridge is: {np.mean(scores_r)}')
print(f'The 5-fold cv mean score of ridge with scaling is: {np.mean(scores_r_scaled)}\n')

print(f'The 5-fold cv mean score of lasso is: {np.mean(scores_l)}')
print(f'The 5-fold cv mean score of lasso with scaling is: {np.mean(scores_l_scaled)}\n')

print(f'The 5-fold cv mean score of elastic net is: {np.mean(scores_e)}')
print(f'The 5-fold cv mean score of elastic net with scaling is: {np.mean(scores_e_scaled)}\n')
```

```
The 5-fold cv mean score of linear regression is: 0.7555894947476223
The 5-fold cv mean score of linear regression with scaling is: 0.7554271073053143
```

```
The 5-fold cv mean score of ridge is: 0.5388338881740073
The 5-fold cv mean score of ridge with scaling is: 0.7567782330762535
```

```
The 5-fold cv mean score of lasso is: 0.7554803973209105
The 5-fold cv mean score of lasso with scaling is: 0.7554798334915229
```

```
The 5-fold cv mean score of elastic net is: 0.593946339788286
The 5-fold cv mean score of elastic net with scaling is: 0.5809007178203834
```

Observation:

For ridge regression model, scaling helps to increase the mean cv score, but for other three models, scaling does not seem to much improve the results.

Part 5: Parameter Tuning

```
In [16]: from sklearn.model_selection import GridSearchCV
```



```

In [17]: # Ridge
param_grid_r = {'ridge__alpha': np.logspace(-3, 4, 8)}
grid_r = GridSearchCV(model_r_scaled, param_grid_r, return_train_score =
True)
grid_r.fit(X_train, y_train)

# Lasso
param_grid_l = {'lasso__alpha': np.logspace(-3, 4, 8)}
grid_l = GridSearchCV(model_l_scaled, param_grid_l, return_train_score =
True)
grid_l.fit(X_train, y_train)

# ElasticNet
param_grid_e = {'elasticnet__alpha': np.logspace(-3, 4, 8),
'elasticnet__l1_ratio': [0.01, .1, .5, .8, .9, .95, .98,
1]}
grid_e = GridSearchCV(model_e_scaled, param_grid_e, return_train_score =
True)
grid_e.fit(X_train, y_train)

print(f'Best parameter of ridge is: {grid_r.best_params_}')
print(f'Best mean cv score of ridge is: {grid_r.best_score_}')

print(f'Best parameter of lasso is: {grid_l.best_params_}')
print(f'Best mean cv score of lasso is: {grid_l.best_score_}')

print(f'Best parameter of elastic net is: {grid_e.best_params_}')
print(f'Best mean cv score of elastic net is: {grid_e.best_score_}')

Best parameter of ridge is: {'ridge__alpha': 1.0}
Best mean cv score of ridge is: 0.7567782330762535
Best parameter of lasso is: {'lasso__alpha': 100.0}
Best mean cv score of lasso is: 0.7557251406146536
Best parameter of elastic net is: {'elasticnet__alpha': 0.01, 'elasticn
et__l1_ratio': 0.9}
Best mean cv score of elastic net is: 0.7576859367405582

```

Observation:

By using grid search, the mean cv score of elastic net model improves, but for the other two models, grid search does not seem to much improve the results.

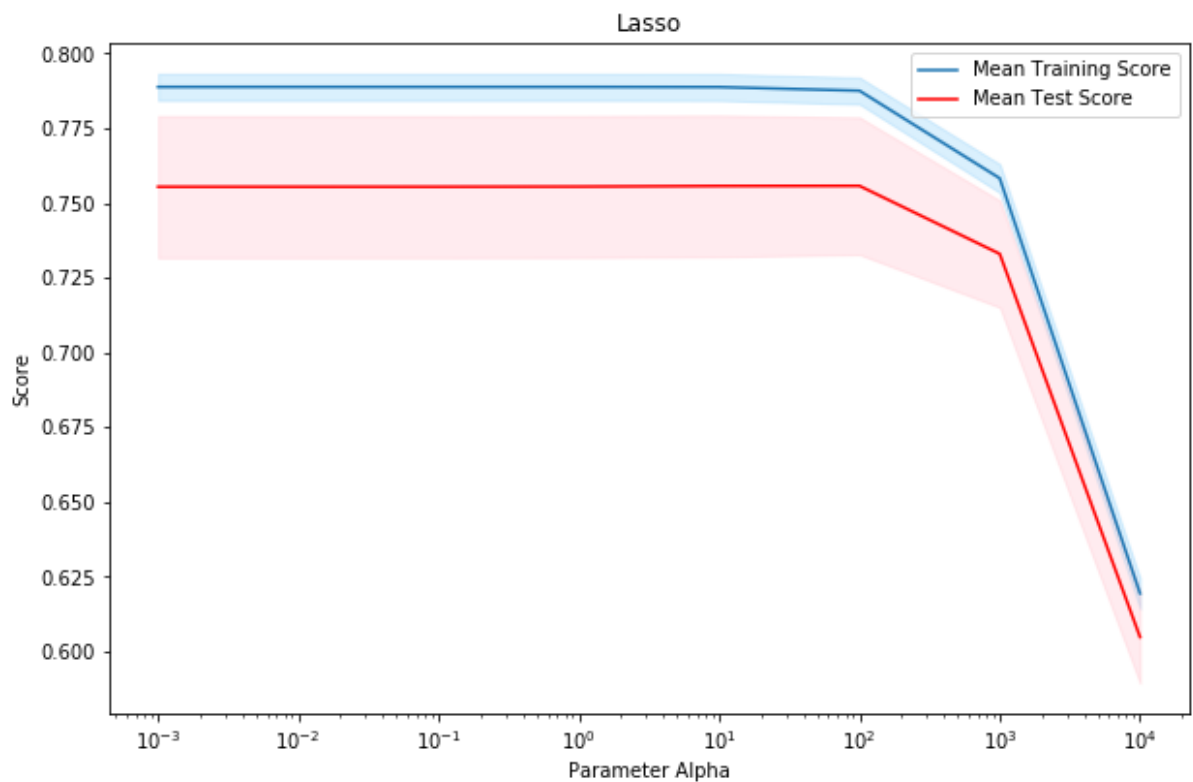
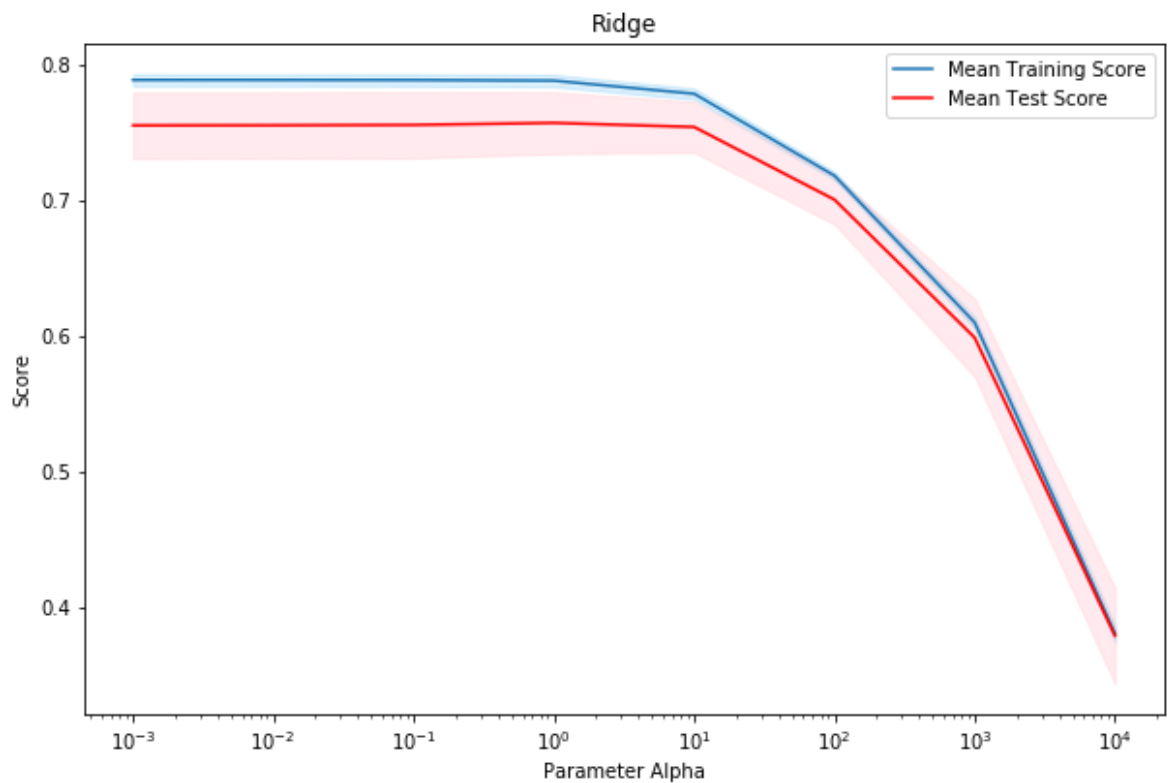
```

In [18]: # Visualization
fig, ax = plt.subplots(2, 1, figsize = (10, 14))
# Ridge
ax[0].plot(np.logspace(-3,4,8),
            grid_r.cv_results_['mean_train_score'],
            label = 'Mean Training Score')
ax[0].fill_between(np.logspace(-3,4,8),
                  grid_r.cv_results_['mean_train_score'] - grid_r.cv_re
sults_['std_train_score'],
                  grid_r.cv_results_['mean_train_score'] + grid_r.cv_re
sults_['std_train_score'],
                  color = 'lightskyblue', alpha = 0.3)
ax[0].plot(np.logspace(-3,4,8),
            grid_r.cv_results_['mean_test_score'],
            color = 'red', label = 'Mean Test Score')
ax[0].fill_between(np.logspace(-3,4,8),
                  grid_r.cv_results_['mean_test_score'] - grid_r.cv_res
ults_['std_test_score'],
                  grid_r.cv_results_['mean_test_score'] + grid_r.cv_res
ults_['std_test_score'],
                  color = 'pink', alpha = 0.3)
ax[0].set_xscale('log')
ax[0].set_xlabel('Parameter Alpha')
ax[0].set_ylabel('Score')
ax[0].set_title('Ridge')
ax[0].legend()

# Lasso
ax[1].plot(np.logspace(-3,4,8),
            grid_l.cv_results_['mean_train_score'],
            label = 'Mean Training Score')
ax[1].fill_between(np.logspace(-3,4,8),
                  grid_l.cv_results_['mean_train_score'] - grid_l.cv_re
sults_['std_train_score'],
                  grid_l.cv_results_['mean_train_score'] + grid_l.cv_re
sults_['std_train_score'],
                  color = 'lightskyblue', alpha = 0.3)
ax[1].plot(np.logspace(-3,4,8),
            grid_l.cv_results_['mean_test_score'],
            color = 'red', label = 'Mean Test Score')
ax[1].fill_between(np.logspace(-3,4,8),
                  grid_l.cv_results_['mean_test_score'] - grid_l.cv_res
ults_['std_test_score'],
                  grid_l.cv_results_['mean_test_score'] + grid_l.cv_res
ults_['std_test_score'],
                  color = 'pink', alpha = 0.3)
ax[1].set_xscale('log')
ax[1].set_xlabel('Parameter Alpha')
ax[1].set_ylabel('Score')
ax[1].set_title('Lasso')
ax[1].legend()

```

Out[18]: <matplotlib.legend.Legend at 0x123d8ff90>



```
In [19]: train_score_e = pd.pivot_table(pd.DataFrame(grid_e.cv_results_),
    values='mean_train_score', index='param_elasticnet_alpha', columns=
    'param_elasticnet_l1_ratio')
test_score_e = pd.pivot_table(pd.DataFrame(grid_e.cv_results_),
    values='mean_test_score', index='param_elasticnet_alpha', columns=
    'param_elasticnet_l1_ratio')
```

```
In [20]: # ElasticNet
fig, ax = plt.subplots(2, 1, figsize = (15, 20))
im = ax[0].imshow(train_score_e, cmap = 'Oranges')
ax[0].set_xticks(np.arange(8))
ax[0].set_yticks(np.arange(8))
ax[0].set_xticklabels([0.01, .1, .5, .8, .9, .95, .98, 1])
ax[0].set_yticklabels(np.logspace(-3,4,8))

for i in range(8):
    for j in range(8):
        text = ax[0].text(j, i, round(train_score_e.iloc[i, j], 4),
                           ha='center', va='center')

ax[0].set_title('Mean Train Score ElasticNet Grid Search')
ax[0].set_xlabel('L1 Ratio')
ax[0].set_ylabel('Alpha')

cbar = ax[0].figure.colorbar(im, ax = ax[0])
cbar.ax.set_ylabel('Mean Train Score', rotation=-90, va="bottom")



im = ax[1].imshow(test_score_e, cmap = 'Oranges')
ax[1].set_xticks(np.arange(8))
ax[1].set_yticks(np.arange(8))
ax[1].set_xticklabels([0.01, .1, .5, .8, .9, .95, .98, 1])
ax[1].set_yticklabels(np.logspace(-3,4,8))


for i in range(8):
    for j in range(8):
        text = ax[1].text(j, i, round(test_score_e.iloc[i, j], 4),
                           ha='center', va='center')

ax[1].set_title('Mean Test Score ElasticNet Grid Search')
ax[1].set_xlabel('L1 Ratio')
ax[1].set_ylabel('Alpha')

cbar = ax[1].figure.colorbar(im, ax = ax[1])
cbar.ax.set_ylabel('Mean Test Score', rotation=-90, va="bottom")
```

```
Out[20]: Text(0, 0.5, 'Mean Test Score')
```





Part 6: 20 Most Important Coefficients

```
In [21]: ohe = OneHotEncoder()
ohe.fit(X_train[X_train.select_dtypes(include=['object']).columns], y_train)
all_columns = np.hstack([X_train.select_dtypes(include=['float64', 'int64']).columns, ohe.get_feature_names()])

# Ridge
index_r = list(np.argsort(np.abs(grid_r.best_estimator_['ridge'].coef_)))[-20:]
estimate_r = [list(grid_r.best_estimator_['ridge'].coef_)[i] for i in index_r]

# Lasso
index_l = list(np.argsort(np.abs(grid_l.best_estimator_['lasso'].coef_)))[-20:]
estimate_l = [list(grid_l.best_estimator_['lasso'].coef_)[i] for i in index_l]


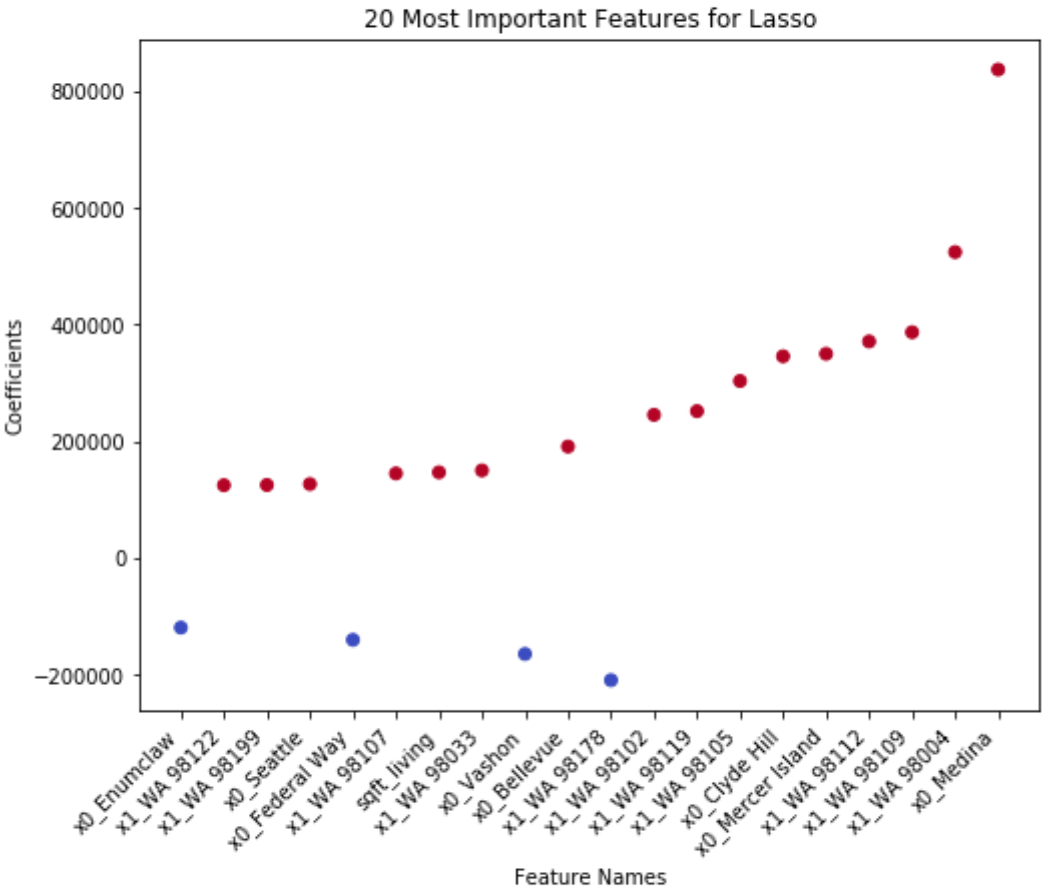
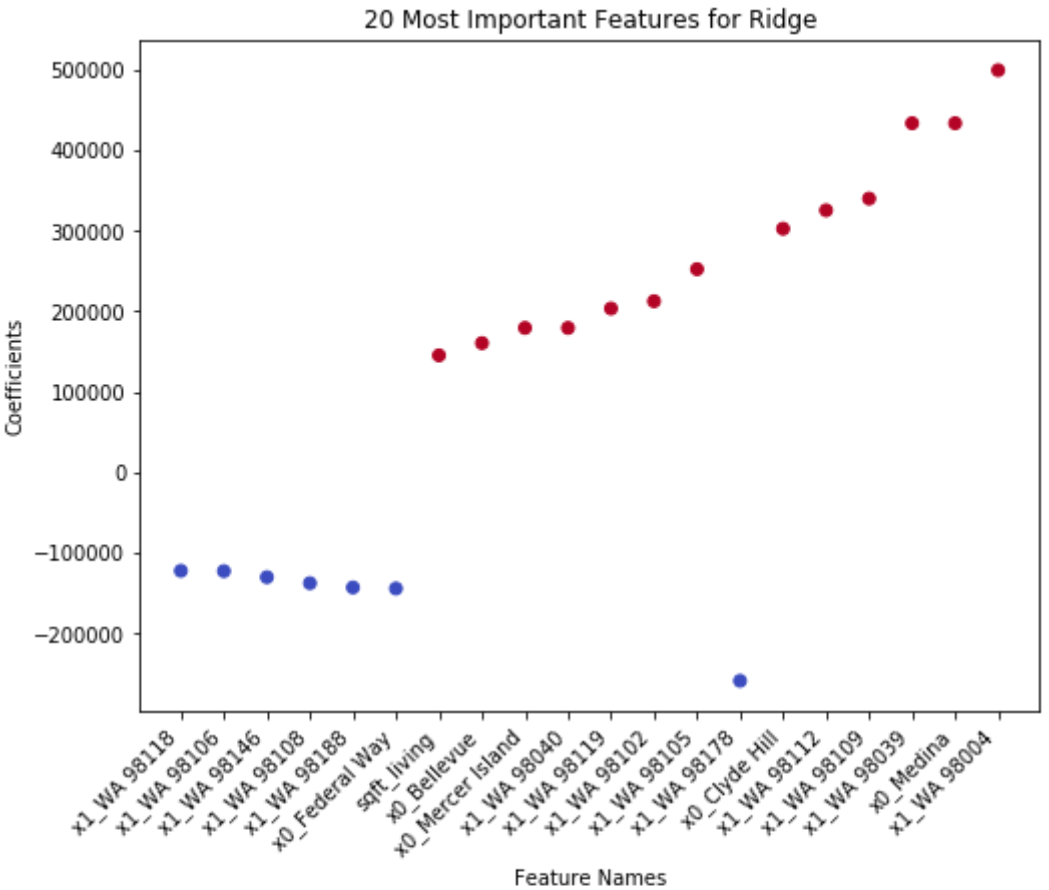
# ElasticNet
index_e = list(np.argsort(np.abs(grid_e.best_estimator_['elasticnet'].coef_)))[-20:]
estimate_e = [list(grid_e.best_estimator_['elasticnet'].coef_)[i] for i in index_e]
```


```
In [22]: # Visualization
fig, ax = plt.subplots(3, 1, figsize = (8, 20))
ax[0].scatter(all_columns[index_r], estimate_r, c = np.sign(estimate_r),
cmap = 'coolwarm')
ax[0].set_ylabel('Coefficients')
ax[0].set_xlabel('Feature Names')
ax[0].set_title('20 Most Important Features for Ridge')
ax[0].set_xticklabels(all_columns[index_r], rotation=45, horizontalalign
ment='right')

ax[1].scatter(all_columns[index_l], estimate_l, c = np.sign(estimate_l),
cmap = 'coolwarm')
ax[1].set_ylabel('Coefficients')
ax[1].set_xlabel('Feature Names')
ax[1].set_title('20 Most Important Features for Lasso')
ax[1].set_xticklabels(all_columns[index_l], rotation=45, horizontalalign
ment='right')

ax[2].scatter(all_columns[index_e], estimate_e, c = np.sign(estimate_e),
cmap = 'coolwarm')
ax[2].set_ylabel('Coefficients')
ax[2].set_xlabel('Feature Names')
ax[2].set_title('20 Most Important Features for Elastic Net')
ax[2].set_xticklabels(all_columns[index_e], rotation=45, horizontalalign
ment='right')

fig.tight_layout(pad=3.0)
```



Observation:

The three models do not agree on the most important 20 features.